

## William Garrett Corder

- **Multiprocessing**

- Within the first for loop `wc.c` forks to start a number of processes equal to `nChildProcesses` called `looping_process`. Each of these processes creates a new `fork()` in `children[]` to run `word_count` using `offset[i]` and `bytes_to_count[i]` such that each process has a designated portion of the file to search.
- `looping_process` felt necessary to include such that all sections of the file could be searched simultaneously.

- **IPC**

- Each child process has a pipe it correlates to stored in `fd[][]`.
- When the process in `children[i]` successfully finishes running `word_count` and collects data in the `count_t` count it writes this data to `fd[i][1]`.
- Once all child processes have finished running another for loop iterates through the pipes and reads the count data from `fd[i][0]`.

- **Error Handling**

- The outer process called `looping_process` handles the errors using the variable `status`. Each time the child process of `looping_process` concludes `status` is updated.
- When `status = 0` it means that the child process exited safely and has written its information to the appropriate pipe.
- If `status` does not equal 0 it means that the child process did not exit safely so it restarts the child process to run again.

### Code Overview

- `wc.c` uses 4 arrays based on the variable, `nChildProcesses`, given in the command line argument. Each value in the arrays corresponds to one of the child processes intended to run `word_count()`.
  - The function `fill_offset_arrays()` calculates the bytes each process will need to search as well as the offset value they will need to start from then places them into the arrays `offsets[]` and `bytes_to_count[]`.
- The bulk of `main()` happens in the first for loop which iterates the variable `i` from 0 to `nChildProcesses - 1`.
  - First it opens the pipe `fd[i]` to transfer data from the process in `children[i]` to the parent process via IPC
  - It creates a “`looping_process`” for each iteration of the for loop. Each `looping_process` is intended to run `word_count` on the subprocess `children[i]` until the function is successful. The purpose of doing this in a subprocess is to search each section of the file concurrently.
    - Opens the file pointed to by command line arguments, each process needs the file to be opened for it to search through.
    - Initializes the variable `status` and enters a while loop until `status = 0` meaning the process reached `exit(0)` and exited safely.
    - Creates another `fork()` called `children[i]` inside of `looping_process` to run `word_count` while the parent of `children[i]` waits for an exit status.
      - `children[i]` runs `word_count()` and either crashes or succeeds. If it succeeds it writes the count data to `fd[i][1]` to be retrieved by the parent of `looping_process` then exits.
    - When `status = 0` `looping_process` exits and the file is closed.
- Once the for loop ends the main function waits for all `children` processes to conclude. Then it enters another for loop to iterate through all of the pipes in `fd[][]` correlating to the `children` processes and collect the count data stored in the pipes.
- Once all the totals are added the results are printed